

A S T A R T E R K I T · 2 0 2 6

Welcome to the World of AI

made for Tezeta — Berkeley, Library Science & AI

Your name means memory — the longing to hold onto what matters. You are about to spend your life learning how humanity keeps, organizes, and protects exactly that. You were practically named for this.

made with love by your cousin ♥

Start here

This is a map, not a textbook. Read it in any order. Skip what's boring. Come back whenever you feel lost. It's not homework — it's a thing to keep.

What's inside

Part	What it's for
1 · A Letter	A personal note — read this first.
2 · Getting Started	What AI actually is, in plain language. No math.
3 · Using AI Well	How to prompt, verify, and not get fooled.
4 · AI for Library Science	Why your field + AI is a superpower.
5 · Ethics & Staying Human	The rules that matter most.
6 · First Projects	Five things to actually try this week.
7 · Glossary	Decoder ring for the buzzwords.
8 · Resources	Where to keep learning.
9 · Cheat Sheet	The whole thing on one printable page.
10 · Dear Future Tezeta	A letter to fill in now, open at graduation.

If you read nothing else: AI is astonishingly good at language and astonishingly confident when it's wrong. Treat it like a brilliant, fast, slightly unreliable intern: great for first drafts and explaining hard things — but you are always the editor, the fact-checker, and the one whose name is on the work. Stay curious, stay skeptical, never outsource your judgment. That's the whole game.

1 · A Letter Before You Start

Hey Tezeta —

So. Berkeley. Full scholarship. Library science and AI. I keep re-reading that sentence because I'm so proud of you I don't totally know what to do with myself.

I have to start with your name, because I don't think you picked your path by accident. Tezeta — memory, the ache to hold onto what matters, the thing we don't want to lose. And here you are about to spend your life learning how humanity keeps, organizes, and protects exactly that. You were practically named for this. Knowledge is just memory that outlives the person who had it, and you're going to learn to be its keeper.

A few things I wish someone had told me when I was starting something big:

You belong there. The scholarship isn't a fluke and it isn't charity. Somebody read your story and bet on you. When the impostor feeling shows up — and it will, probably around week three — remember that everyone feels it, and that feeling it means you're in a room worth being in.

AI is a tool, not an oracle. It will sound confident and be wrong. It will write you a beautiful paragraph with a fake citation in it. Your job — your whole training, really — is to be the person who can tell the difference. That makes you more valuable in the AI age, not less.

Stay a beginner on purpose. The field moves so fast that everyone is a beginner at something. Being comfortable not knowing is a superpower.

Don't let the machine think for you — let it think with you. Use AI to go faster and understand hard things. Never use it to skip the part where you understand. The understanding is the thing you're actually there to get.

I made you this little package — guides, guardrails, things to try. Open it when you're curious or stuck, and close it when you've got it.

Go be brilliant, Tezeta. Call me when you find something cool. I want to hear all of it.

Love always,
your cousin ♥

P.S. — This is in my words. If you're printing it, make it yours: tell her the story behind her name, the thing only you'd say, the memory only the two of you share. That's the part no AI could ever write — which is sort of the whole point.

2 · What AI Actually Is

No math. No jargon you don't need. Just the mental model.

The one-sentence version

Modern AI (the chatbot kind) is a system that has read an enormous amount of human writing and learned to predict what words come next — so well that "predicting the next word" starts to look like understanding, reasoning, and conversation. That's it. Everything else is detail.

A useful picture

Imagine someone who has read a huge slice of the internet, millions of books, and mountains of code — but who has **no memory of any single source, no ability to look things up** (unless you give it tools), and **no inner sense of "I'm not sure."** That's a large language model (LLM): the engine behind Claude, ChatGPT, and Gemini. It explains the two things that confuse people most:

- **Why it's so good:** it has seen how experts write about nearly everything.
- **Why it makes things up ("hallucinates"):** it predicts plausible-sounding text, not verified facts. A fake citation looks exactly like a real one to the model.

What it's genuinely great at

- Explaining hard ideas at any level ("explain like I'm 12 / like a grad student")
- First drafts of almost anything — then you edit
- Summarizing long things; translating tone, language, and format
- Brainstorming — "give me 20 angles, then I'll pick"
- Rubber-ducking: talking through a problem until you see the answer

What it's bad at (know these cold)

- **Facts, names, numbers, citations** — verify every one
- **Recent events** past its training cutoff (unless it has live search)
- **Math and counting** without a tool to help
- **Knowing what it doesn't know** — it rarely says "I'm not sure"

Trust nothing factual until you've checked it. Use AI for thinking; use real sources for truth. You, a future librarian, are being trained to be exactly the person who knows how to check. That's why this field is about to matter more than ever.

3 · How to Use AI Well

Not secret tricks — just a handful of habits.

1. Write better prompts — the C.R.A.F.T. checklist

- **Context** — who you are, what's going on
- **Role** — who you want the AI to be ("act as a patient research librarian")
- **Ask** — the actual task, specific
- **Format** — how you want the answer
- **Tone** — length and style

Weak: "Tell me about AI in libraries."

Strong: "Act as a research librarian. I'm a first-year student writing a 5-page paper. Give me three specific, debatable research questions about AI in public libraries, each with one reason it matters. Bullet points, plain English."

2. Iterate — don't accept the first answer

"Make it simpler." "What's the strongest counterargument?" "Give me three more, weirder this time." The good stuff usually shows up in round two or three.

3. Verify everything factual (the 30-second rule)

If a fact, name, date, statistic, or citation matters, confirm it from a real source first. Especially **citations** — AI invents real-looking sources with completely fake details, and your professors will catch it. Library databases (JSTOR, your catalog, Google Scholar) are your truth layer.

4. Never paste in things you shouldn't

Treat a chatbot like a postcard, not a diary. No passwords, no other people's private data, no confidential work, no unpublished work you're not ready to risk.

5. Keep your fingerprints on the work

Draft with AI, then rewrite in your own voice and make sure you understand every line. If you couldn't explain it to a friend with the chatbot closed, you're not done.

6. The trust thermometer

Stakes	How much to verify
	Low — sanity-check

Brainstorming, explaining a concept to yourself

A draft you'll heavily edit

Medium — verify key facts

Anything you'll submit, cite, or publish

High — verify everything, cite real sources

Health, legal, financial, safety decisions

Don't rely on it — ask a qualified human

You're the pilot. AI is the autopilot. Autopilot is amazing — and you never take your hands fully off the controls.

4 · AI for Library Science

A secret not enough people have noticed: library science and AI are the same project wearing different clothes.

Both answer one question — how do humans find the right knowledge at the right moment and trust it? Librarians have built the findable systems forever: catalogs, classification, metadata, reference. AI is a new, powerful way to do information retrieval and organization. You're not learning a field AI threatens — you're learning the field that makes sense of AI.

Concepts where your two worlds collide

- **Information retrieval** — the science of finding relevant things; the shared root of search, databases, and AI.
- **Metadata** — data about data. AI consumes it to work better and can help generate it at scale.
- **Classification & taxonomies** — Dewey, LCSH, ontologies. AI suggests; humans decide what's fair and unbiased.
- **RAG (Retrieval-Augmented Generation)** — the one to know. AI that looks things up in a trusted collection before answering. Basically "give the chatbot a library card."
- **Digital archives & preservation** — AI helps with transcription (handwriting!), translation, and search across huge collections.
- **Information literacy** — teaching people to evaluate sources. In the AI age, the most important thing a librarian gives the public.

Where you, the human, stay essential

- **Truth & verification** — was that source real? Does it say what's claimed?
- **Bias & fairness** — AI learns society's biases; someone has to notice and correct them.
- **Privacy** — libraries fiercely protect what people read; AI tools often want exactly that data. You'll be a guardian.
- **Judgment & ethics** — deciding what should be done, not just what can.

The card catalog, the search engine, and the AI assistant are the same idea a century apart: a bridge between a person and the knowledge they're looking for. You're learning to build the next bridge — and to make sure it's honest, open, and for everyone.

5 · Ethics & Staying Human

The technical stuff you'll pick up fast. This is the part that actually matters.

The line that keeps you out of trouble: Use AI to do your work better. Never use it to avoid doing your work. The honest uses are the ones that leave you smarter at the end.

Academic integrity, concretely

- **Ask first.** Each professor and assignment may differ. "Is AI allowed here, and for what?"
- **Disclose when unsure.** "I used AI to brainstorm and check grammar" is almost always the safe move.
- **Never submit unverified AI text as fact.** Fake citations are the classic way students get caught.
- **The work has to be yours.** If you can't explain it without the chatbot open, you haven't learned it.

Things AI gets wrong — watch for them

- **Bias.** Models inherit human prejudice. Noticing and naming it is part of your job.
- **Hallucination.** Confident, fluent, and wrong. Assume facts need checking.
- **Erasure.** AI can sideline smaller voices, dialects, and histories. Libraries exist partly to protect those.
- **False neutrality.** "The computer said so" feels objective. It isn't.

Privacy — your field's superpower value

Libraries hold one of the strongest privacy ethics on earth: what you read is nobody's business. Many AI tools are built on the opposite instinct. Start the habit now — don't feed private data into chatbots, and champion the patron's right to read without being tracked.

Staying human

Keep your curiosity, your voice, your relationships, and your struggle. The hard, slow, confusing part of learning is the learning — don't let AI smooth it all away.

Let AI make you faster, never make you smaller. Stay curious, stay honest, stay kind — the point of every tool is to help real people, which is the same reason libraries exist.

6 · Your First Five Projects

Reading about AI is like reading about swimming. Get in the water. Pick a free assistant (Claude, ChatGPT, Gemini all have free tiers) and go.

1 · "Explain it five ways"

Take something confusing. Ask the AI to explain it like you're 10, like a grad student, with a cooking analogy, in 3 bullets, and as a short story. Feel how flexible it is.

2 · Catch a hallucination on purpose

Ask for five scholarly sources on the history of library classification, with authors and years — then **check every one** in Google Scholar. Some will be invented. This is the most important lesson here, and you only believe it once you've caught it yourself.

3 · Build a research plan

"Turn this broad topic into three debatable questions." "What sources would I need?" "What's the strongest argument against my position?" "What am I not thinking of?" You did the thinking; it widened your view.

4 · Compare three models

Ask Claude, ChatGPT, and Gemini the same meaty question. Where do they agree? Which was most honest about uncertainty? You'll develop taste for which to reach for.

5 · Make AI tutor you, then quiz you

"Teach me this step by step, checking my understanding after each step." "Now quiz me, one question at a time." A patient tutor at 2 a.m. that never sighs at your questions. Use it a lot.

Every time you use AI, ask: did this make me understand more, or let me understand less? Aim for "more," every time.

7 · Glossary

Plain-English definitions. Come here when a buzzword trips you up.

Term	What it means
LLM (Large Language Model)	AI trained on huge amounts of text that predicts likely next words — the engine behind chatbots.
Generative AI	AI that creates new content (text, images, audio), not just sorts existing things.
Prompt	What you type in: your question, instruction, and context.
Token	A chunk of text, ~3/4 of a word. Models read and write in tokens.
Context window	The model's short-term memory — how much it can hold in mind at once.
Hallucination	When AI states something false with total confidence. The word to remember.
Training cutoff	The date after which the model wasn't trained, so it doesn't natively know recent events.
Bias	Unfair patterns absorbed from training data, reflecting real-world prejudice.
RAG	Retrieval-Augmented Generation — AI that looks things up in a trusted source before answering. "A chatbot with a library card."
Semantic search	Search by meaning rather than exact keywords.
Embedding	Turning text into numbers that capture meaning, so a computer can measure similarity. Engine behind semantic search and RAG.
Metadata	Data about data (author, date, subject). The lifeblood of catalogs.
Multimodal	Handles more than text: images, audio, sometimes video.
Agent	An AI set up to take actions and use tools on your behalf, not just chat.

Don't memorize this. Skim once, return when needed. Understanding beats vocabulary.

8 · Where to Keep Learning

Curated, not overwhelming. Pick one thing that sparks you and follow the thread. Treat specific tools as "true as of 2026" — the ideas age slower than the products.

Start here (gentle, free)

- **Just talk to the tools.** Free tiers of Claude, ChatGPT, and Gemini. Two weeks of daily use teaches more than any course.
- **Elements of AI** — a free, famous, non-technical intro course.
- **Your university library's own guides.** Berkeley will have research guides on AI tools and integrity. It's literally your field.

For the library + AI crossover

- **The American Library Association (ALA)** — their work on AI, privacy, and intellectual freedom is your professional compass.
- Search "**AI literacy for librarians**" — a fast-growing area aimed exactly at you.
- The **ACRL Framework for Information Literacy** — the backbone of evaluating any source.

Books to grow into

- Weapons of Math Destruction — Cathy O'Neil (how algorithms can be unfair)
- Atlas of AI — Kate Crawford (the real-world costs and politics of AI)
- The Alignment Problem — Brian Christian (making AI do what we actually want)

Habits that beat any reading list

- Follow thoughtful people who say "it depends" and "I'm not sure" — the honest ones.
- Keep an **ai-notebook**: prompts that worked, answers that surprised you, times it was wrong.
- Teach someone else — the fastest way to find out what you actually understand.

You do not need to keep up with every headline. Check in monthly, not hourly. Depth beats novelty. And call me when you build something cool. I mean that. ♥

9 · The One-Page Cheat Sheet

Print this. Pin it above your desk. The whole package in 60 seconds.

The golden rule: AI makes you faster. Never let it make you smaller. Use it to think with, never to skip the thinking.

The 5-part prompt — C.R.A.F.T.

Context · Role · Ask · Format · Tone

Before you trust it

- Verify every fact, name, date, and citation. AI invents real-looking sources.
- Ask: does this make me understand more, or less? Aim for more.
- If you can't explain it without the chatbot open, you're not done.

The trust thermometer

Stakes	Verify how much?
Brainstorming / explaining to yourself	Sanity-check
A draft you'll edit	Key facts
Anything you'll submit or cite	Everything — real sources
Health / legal / money / safety	Ask a qualified human

Never paste in

Passwords · other people's private data · confidential work · anything you'd put on a diary page, not a postcard.

In class

Ask **before** each assignment: "Is AI allowed here, and for what?" When unsure, disclose how you used it.

Your superpower: you're training to be the person who can tell true from plausible, real source from fake, fair from biased. In the AI age, that's worth more, not less. ♥

10 · Dear Future Tezeta

A page to fill in by hand — today, before classes start. Seal it. Open it the week you graduate.

Today's date: _____

Right now, I am feeling _____

(nervous? electric? both? write the truth.)

The thing I'm most excited to learn is _____

The thing I'm most scared of is _____

(name it — it gets smaller once it's on paper.)

One reason I belong at Berkeley, even on the days I doubt it:

What I hope AI helps me do (without doing it for me):

A promise to myself about how I'll use these tools honestly:

The kind of librarian I want to become:

Something I never want to lose, no matter how the tech changes:

By the time you read this again, you'll have caught a hundred hallucinations, written things you're proud of, and helped real people find what they were looking for.

Your name means memory. So remember this beginning — the nerves, the hope, all of it. You did something brave today. I'm proud of you. I was always going to be.

— your cousin, and now also you ♥